

Deep Learning on Edge Device

or

Pain and suffering

Why?

- IoT is the “thing”
- Avoid clogging up networks
- Reduce network costs
- Edge devices have increasing GPU capabilities

The Jetsons

- Jetson Nano (~\$240)
 - GPU: 128 core Maxwell
 - Quad core ARM A57
 - 4GB RAM (shared)
 - 2x CSI
 - Gigabit Ethernet
 - HDMI/Display port
 - 4x USB 3
 - 40pin GPIO
 - Power: USB micro or barrel jack (5-10W)



Does it work?

- Kind of...
 - JetPacks (ie prebuilt images) only come with Python 3.6, based on Linux4Tegra (L4T)
 - Compilation is terribly slow (missing wheel files)
 - PyTorch can easily run out of memory
 - Nvidia-based code/models does work

The Jetsons' bigger brother

- Jetson Xavier NX (>\$400)
 - 512 core Volta
 - 8-core ARMv8.2
 - 2x CSI
 - Gigabit Ethernet, M.2
 - HDMI/Display port
 - 4x USB 3
 - 40pin GPIO
 - Power: 19V barrel jack (10-15W)



Does it work?

- Much better
 - still limited with JetPacks and Python 3.6
 - Compilation is way faster
 - Double amount of RAM helps
 - PyTorch works reasonable